

Apache Flink® Training

Intro



Apache Flink® Training

dataArtisans

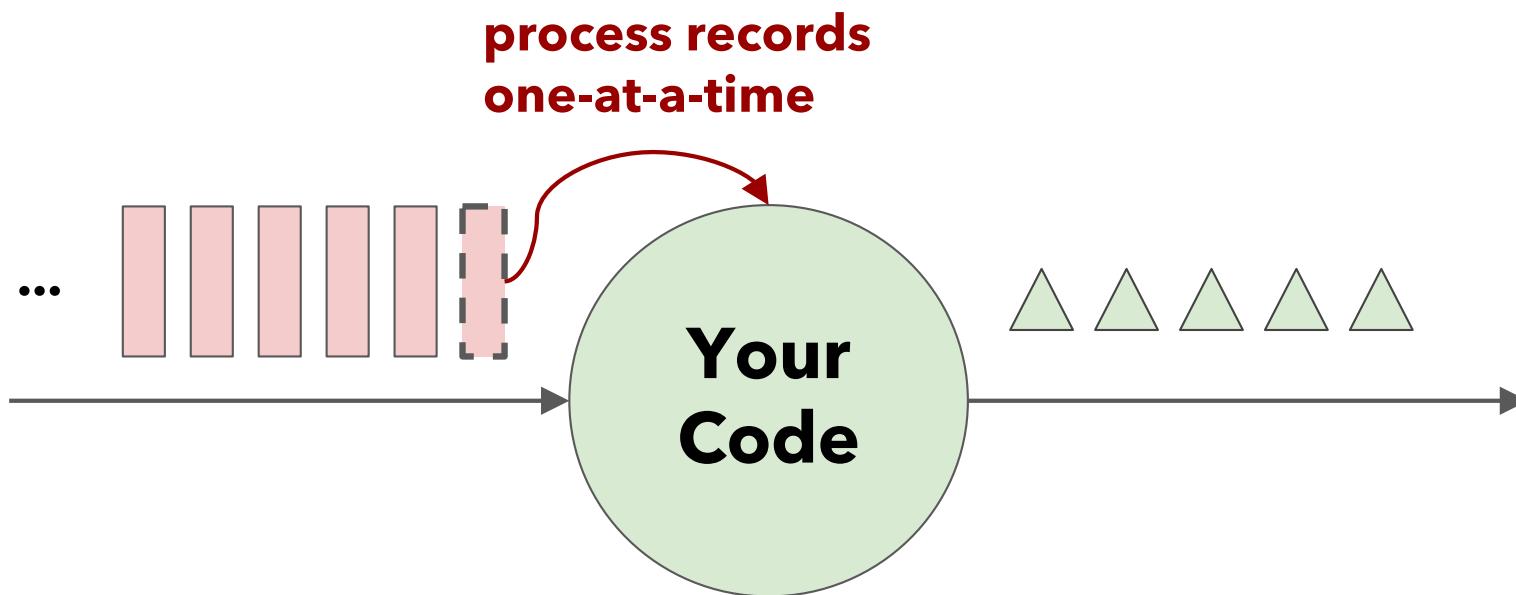
Flink v1.3 – 20.06.2017

Where we're going today



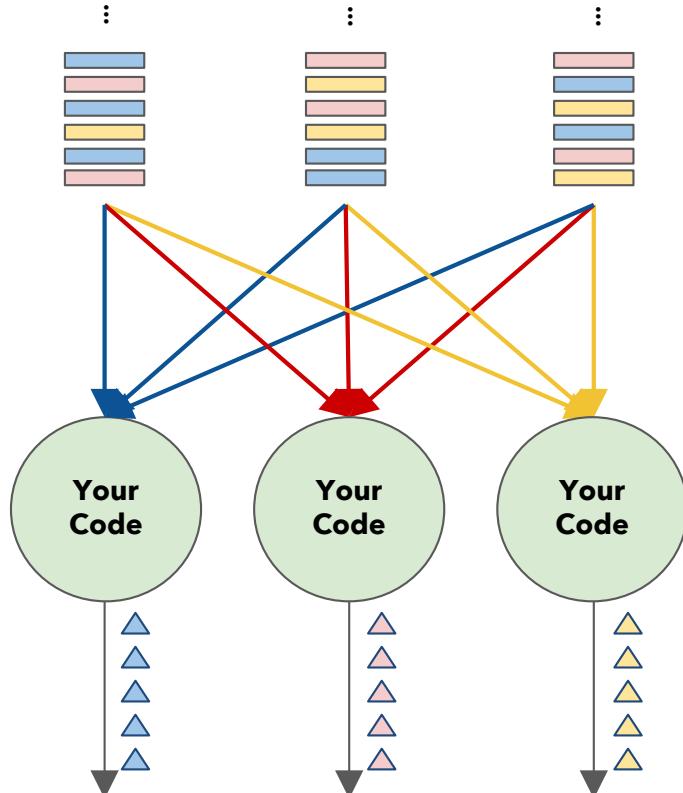
- Stateful stream processing as a paradigm for continuous data
- Apache Flink is a sophisticated and battle-tested stateful stream processor with a comprehensive set of features
- Efficiency, management, and operational issues for state are taken very seriously

Stream Processing



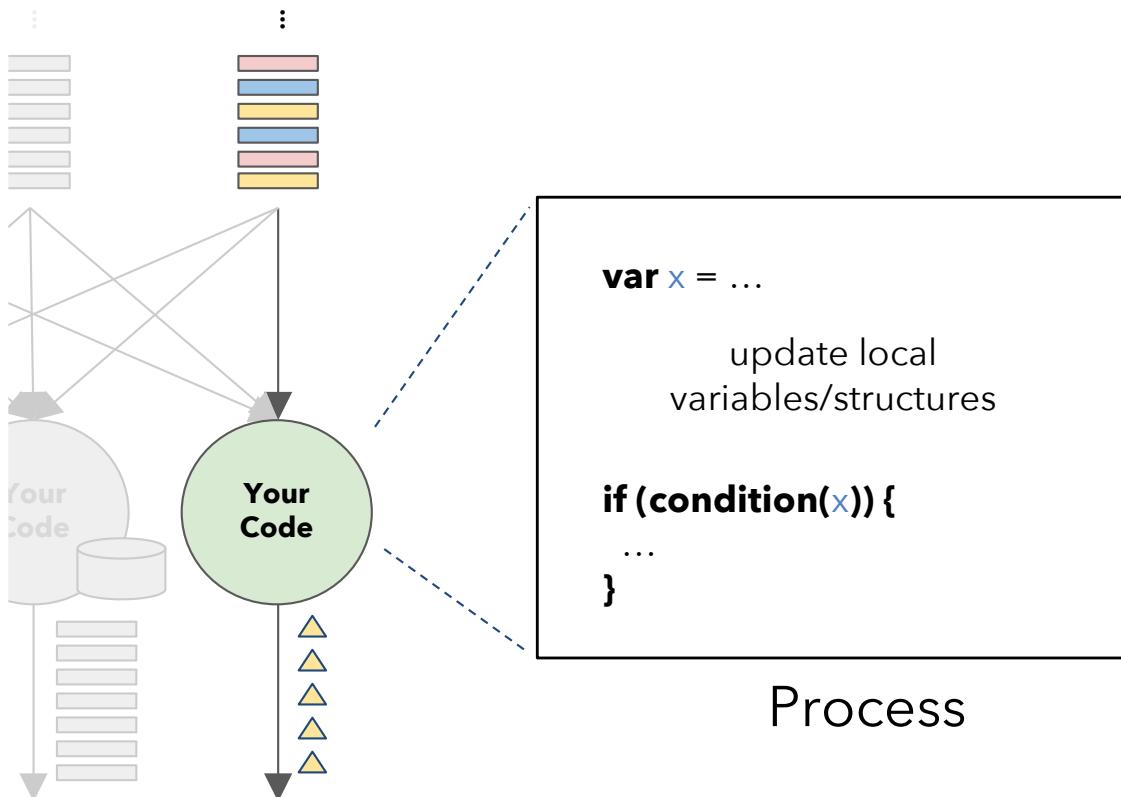
Long running computation, on an endless stream of input

Distributed Stream Processing

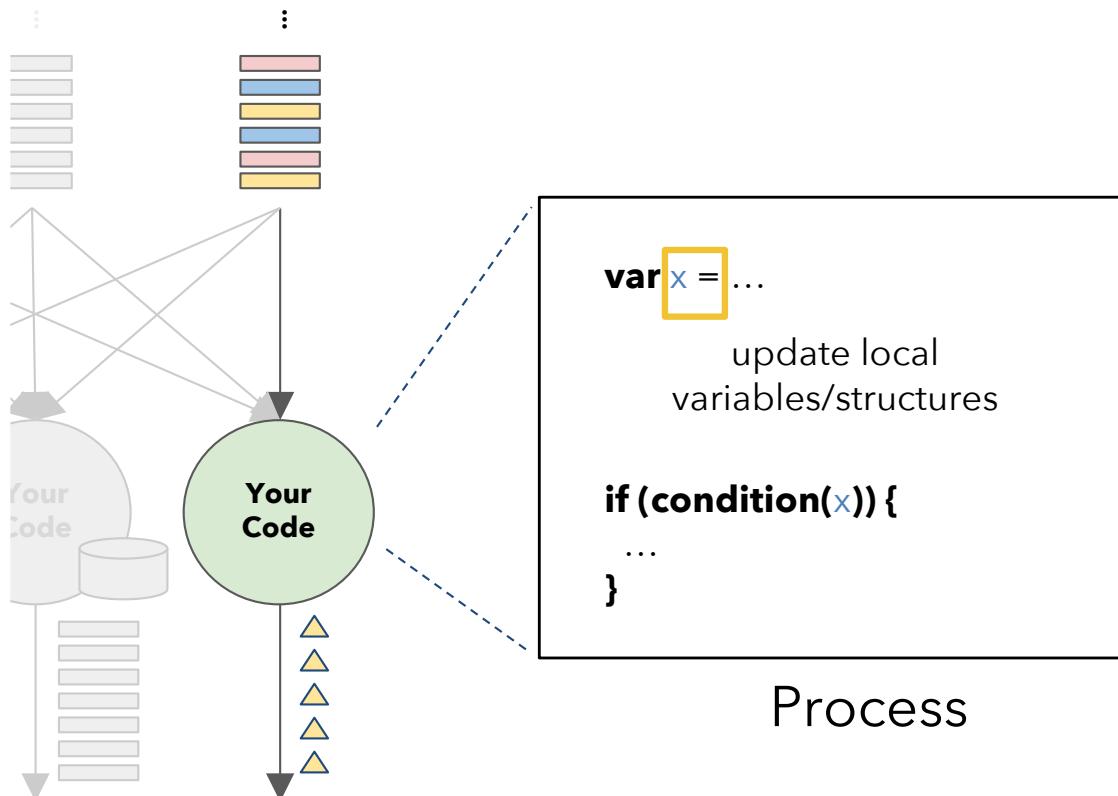


- partitions input streams by some key in the data
- distributes computation across multiple instances
- each instance is responsible for some key range

Stateful Stream Processing

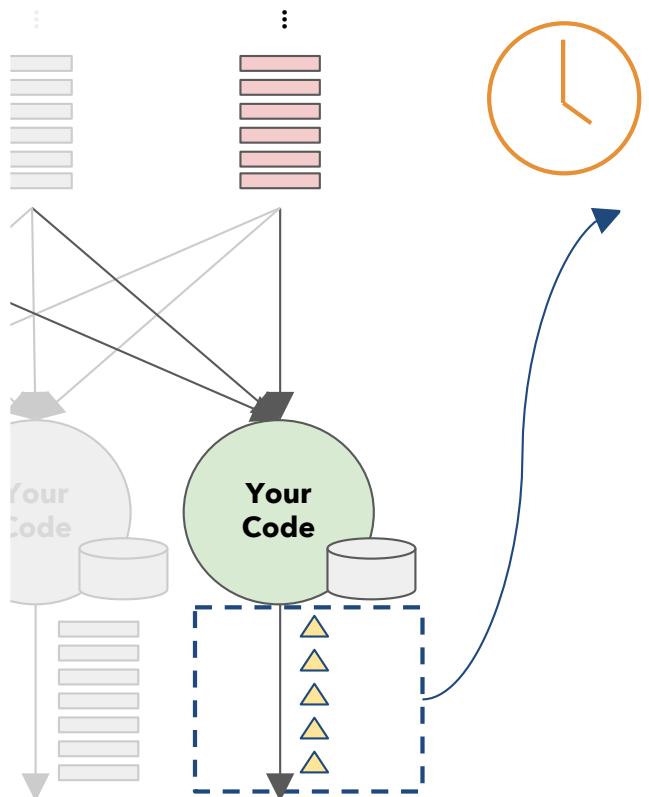


Stateful Stream Processing



- embedded local state backend
- state co-partitioned with the input stream by key

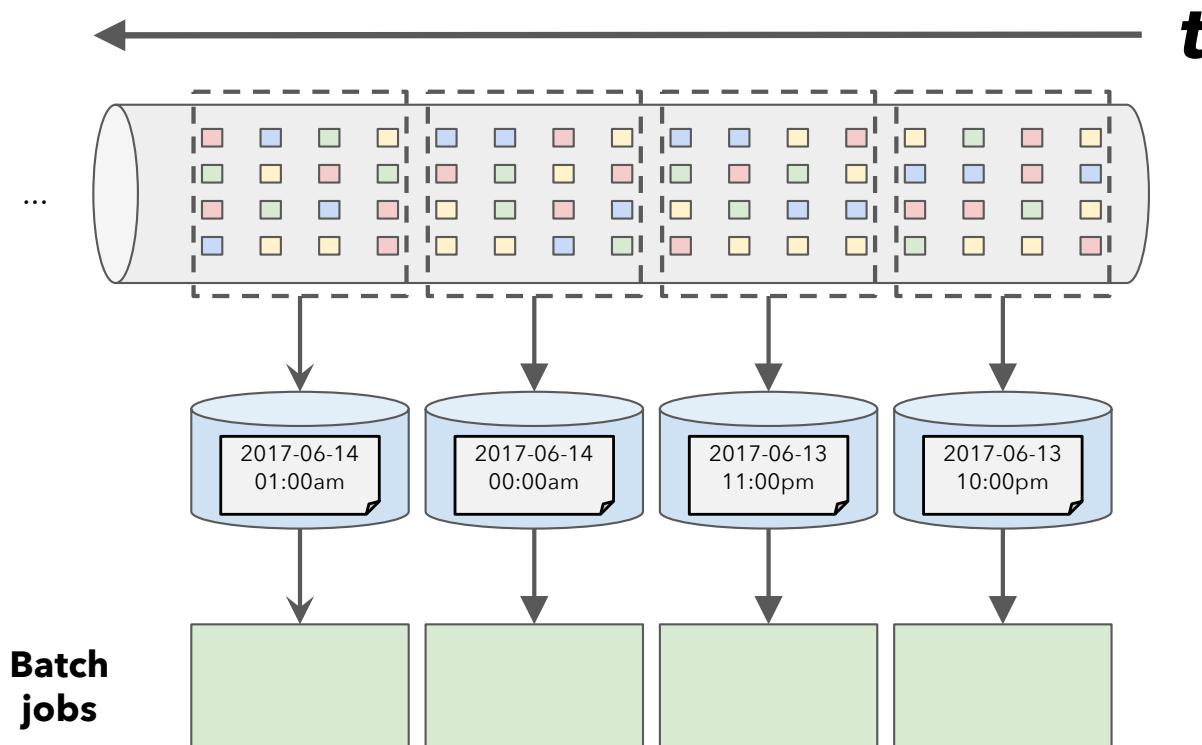
About time ...



When should results be emitted?

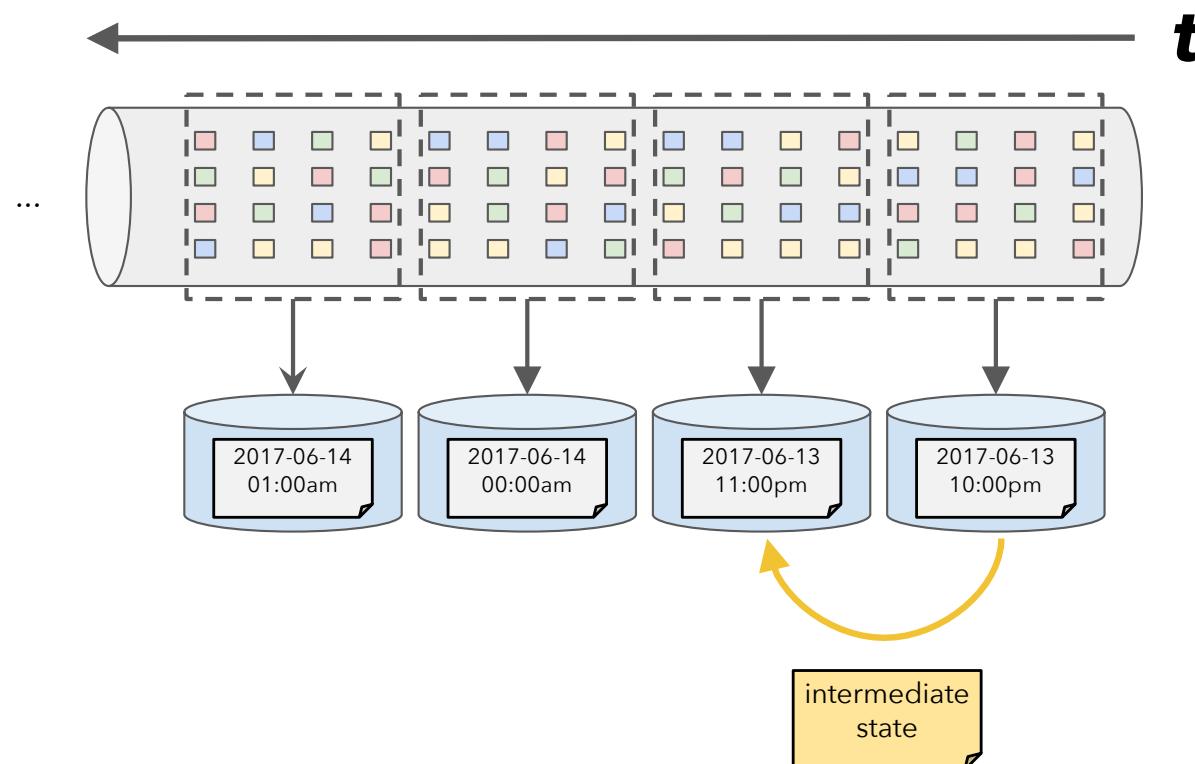
- Control for determining when the computation has fully processed all required events
- It's mostly about time. e.g. have I received all events for 3 - 4 pm?
- Did event *B* occur within 5 minutes of event *A*?
- **Wall clock** time is not correct. **Event-time awareness** is required.

Traditional batch processing



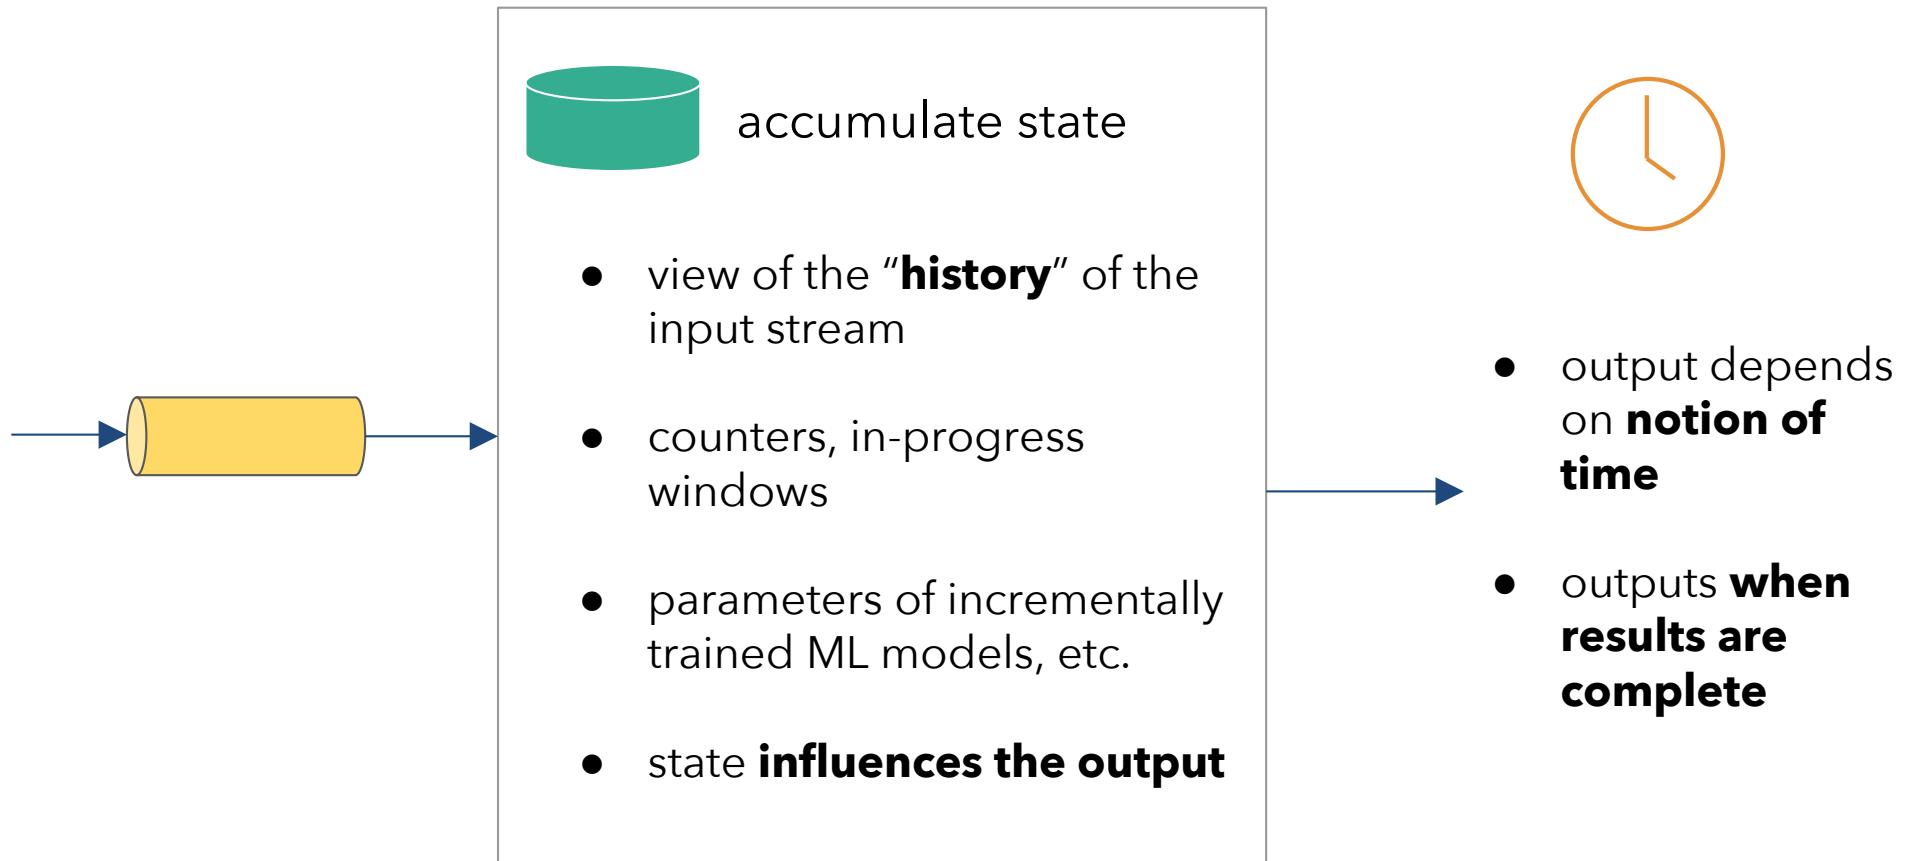
- Continuously ingesting data
- Time-bounded batch files
- Periodic batch jobs

Traditional batch processing (II)

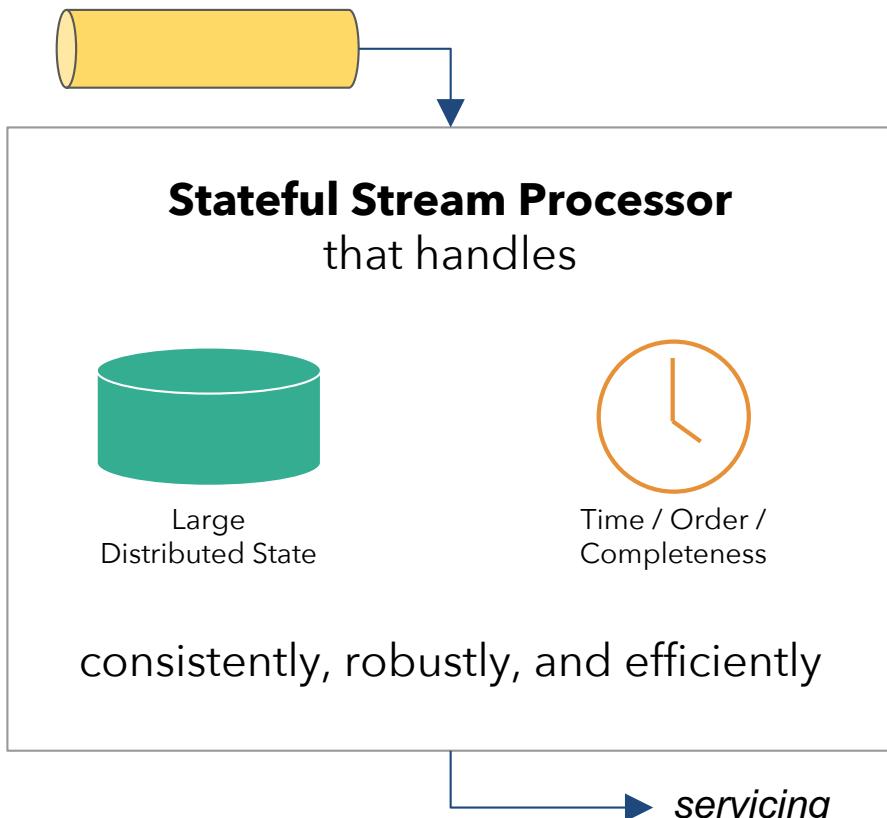


- Consider computing conversion metric (**# of A → B per hour**)
- What if the conversion crossed time boundaries?
→ carry intermediate results to next batch
- What if events come out of order?
→ ???

The ideal way



The ideal way (II)

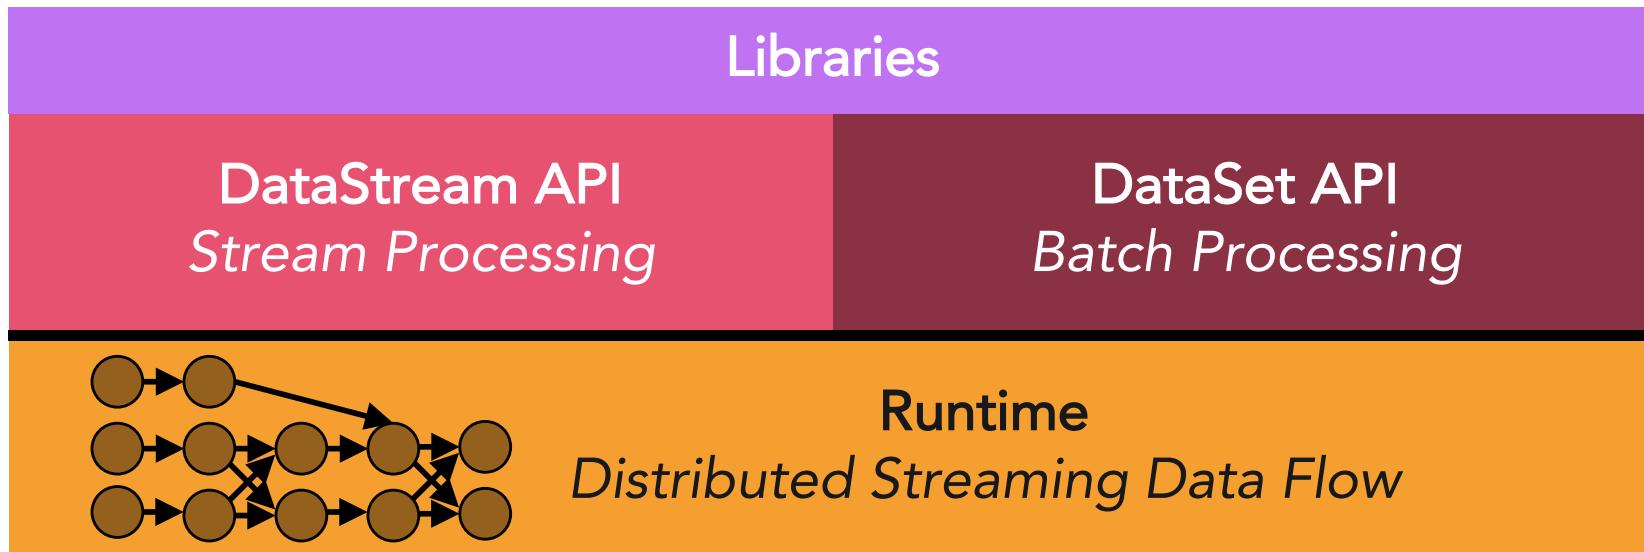


- Stateful stream processing as a **new paradigm** to **continuously process continuous data**
- Produce **accurate** results
- Having results available in real-time (with low latency and high throughput) is a natural consequence of the model
- Process both real-time and historic data using exactly the same application

Flink APIs and Runtime

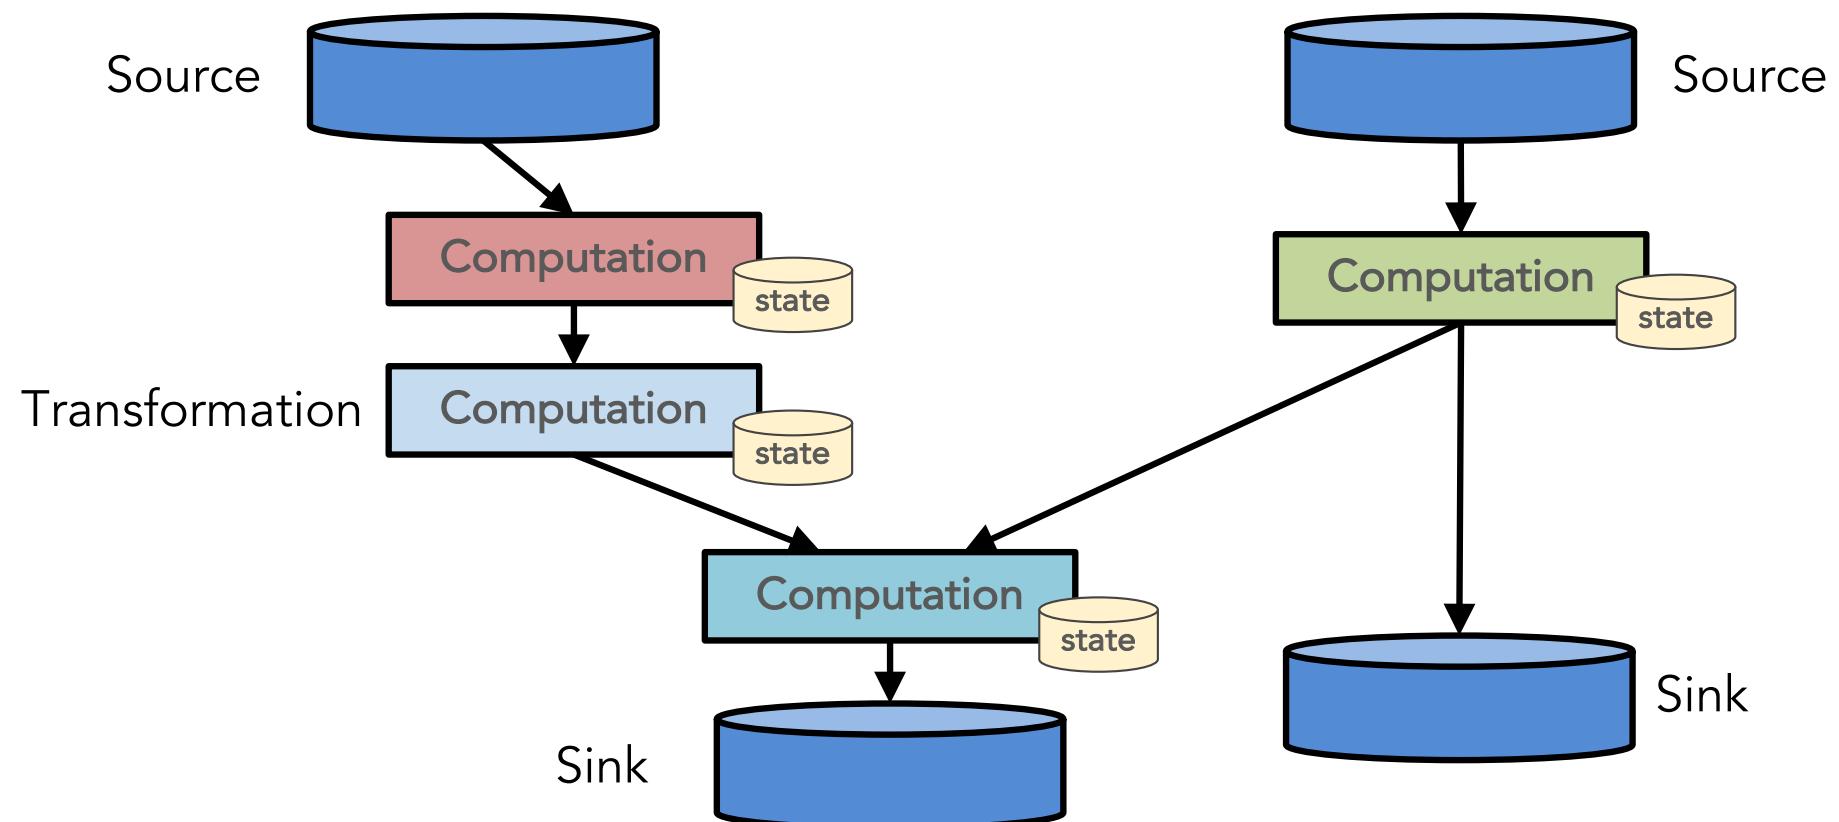


Apache Flink Stack

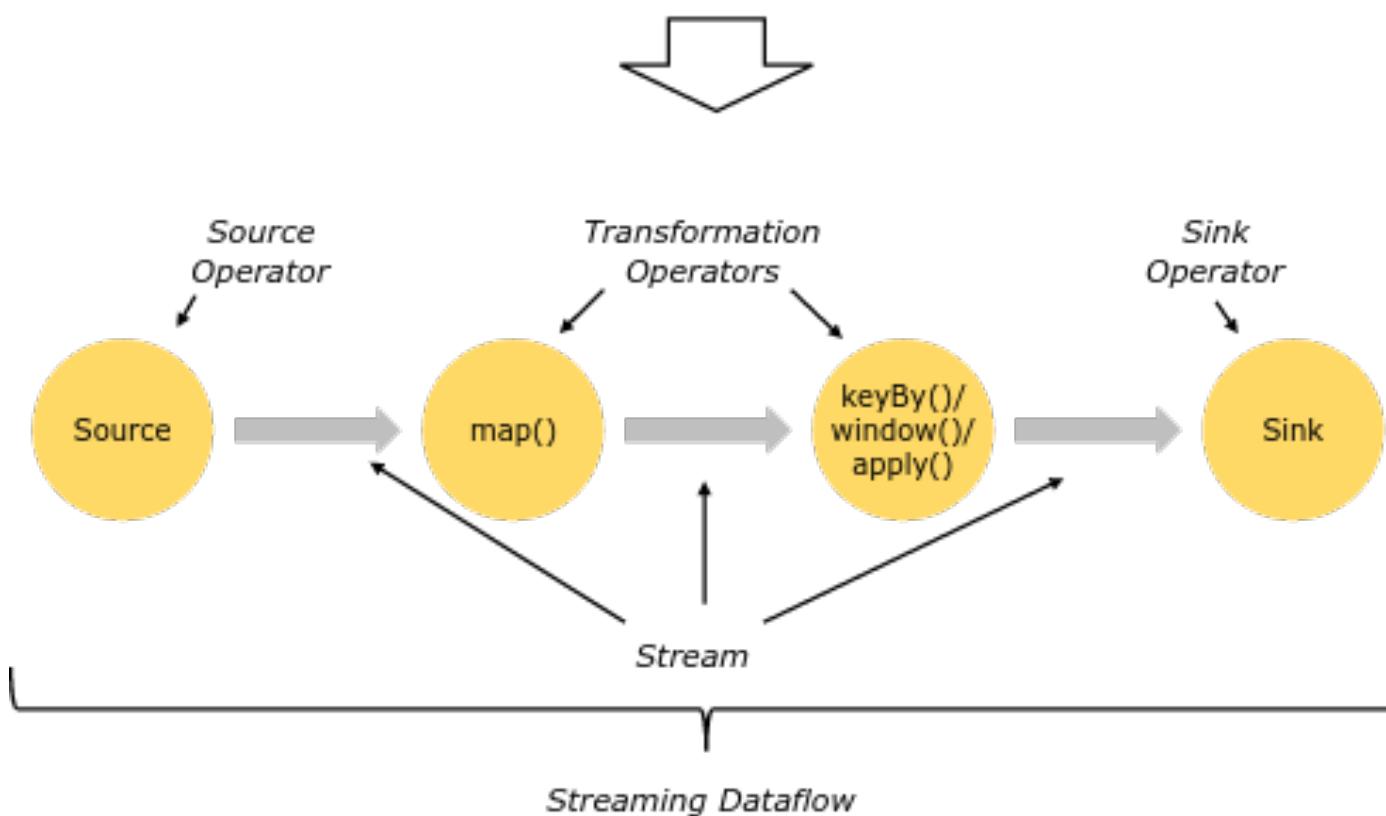


Streaming and batch as first class citizens.

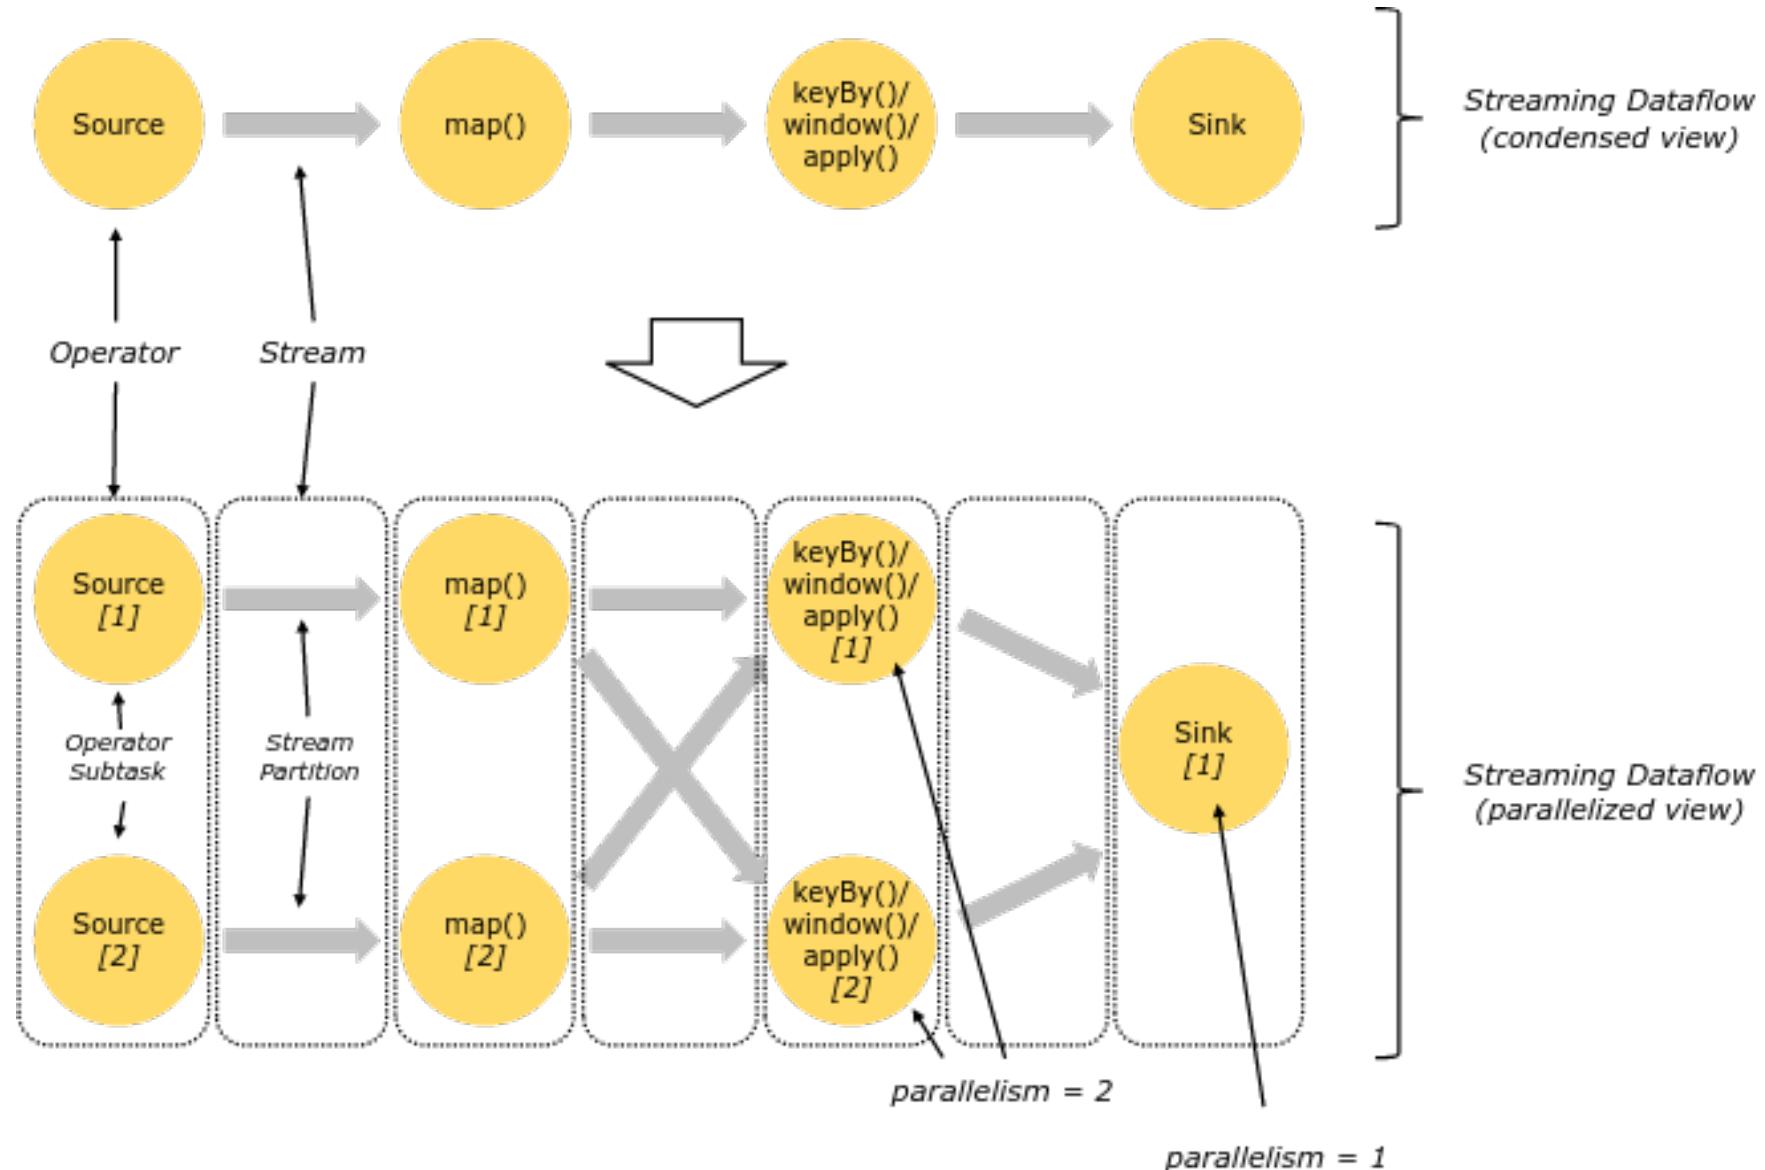
Programming Model



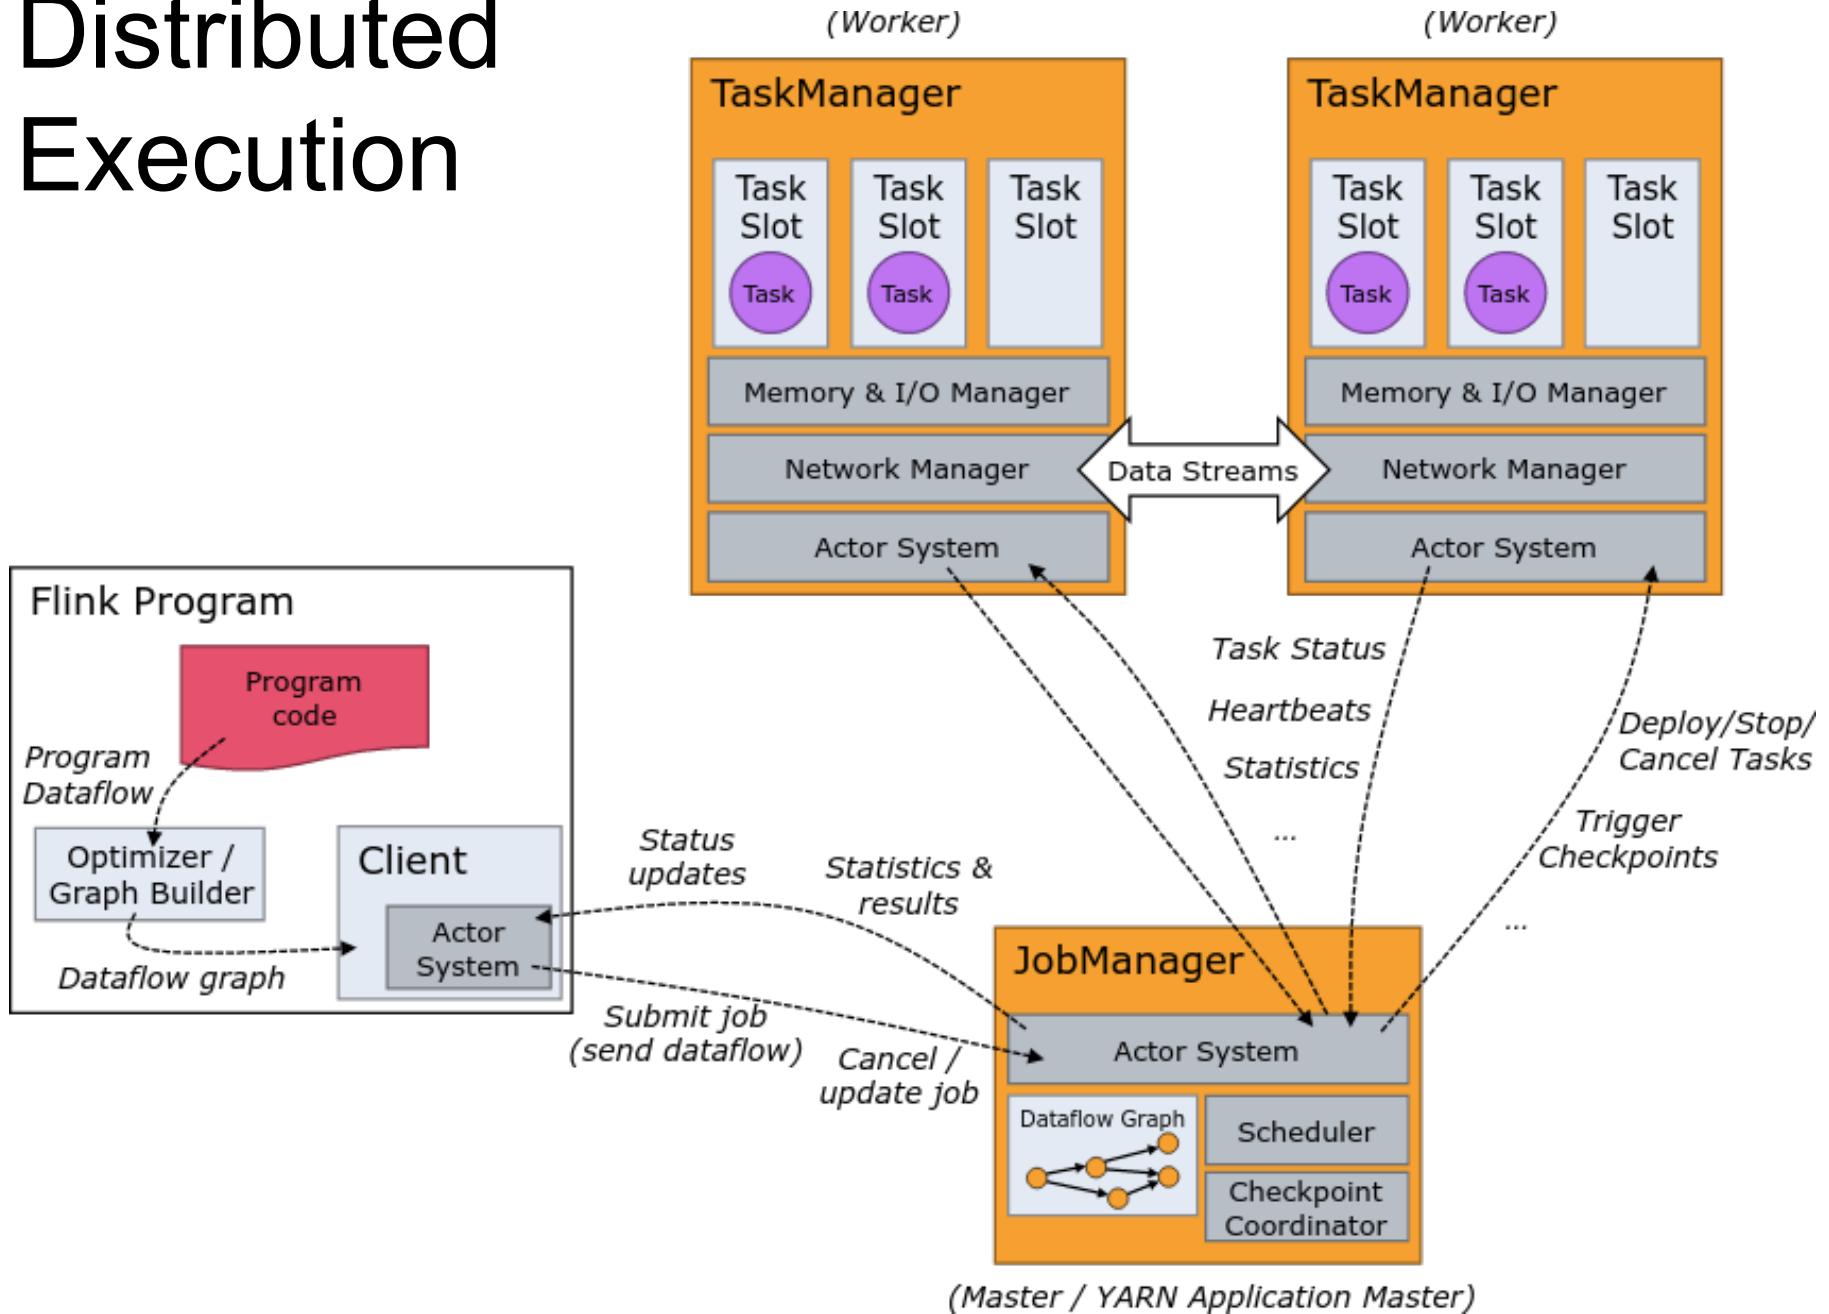
```
DataStream<String> lines = env.addSource(  
    new FlinkKafkaConsumer<>(...)); } Source  
  
DataStream<Event> events = lines.map((line) -> parse(line)); } Transformation  
  
DataStream<Statistics> stats = events  
    .keyBy("id")  
    .timeWindow(Time.seconds(10))  
    .apply(new MyWindowAggregationFunction()); } Transformation  
  
stats.addSink(new RollingSink(path)); } Sink
```



Parallelism



Distributed Execution



Levels of abstraction



Stream SQL

← high-level language

Table API (*dynamic tables*)

← declarative DSL

DataStream API (*streams, windows*)

← stream processing & analytics

Process Function (*events, state, time*)

← low-level
(stateful stream processing)

Process Function



```
class MyFunction extends ProcessFunction[MyEvent, Result] {

    // declare state to use in the program
    lazy val state: ValueState[CountWithTimestamp] = getRuntimeContext().getState(...)

    def processElement(event: MyEvent, ctx: Context, out: Collector[Result]): Unit = {
        // work with event and state
        (event, state.value) match { ... }

        out.collect(...) // emit events
        state.update(...) // modify state

        // schedule a timer callback
        ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
    }

    def onTimer(timestamp: Long, ctx: OnTimerContext, out: Collector[Result]): Unit = {
        // handle callback when event-/processing- time instant is reached
    }
}
```

Data Stream API



```
val lines: DataStream[String] = env.addSource(  
    new FlinkKafkaConsumer10<>(...))  
  
val events: DataStream[Event] = lines.map((line) => parse(line))  
  
val stats: DataStream[Statistic] = stream  
    .keyBy("sensor")  
    .timeWindow(Time.seconds(5))  
    .sum(new MyAggregationFunction())  
  
stats.addSink(new RollingSink(path))
```

Table API & Stream SQL



```
// Table API
val tapiResult: Table = tEnv.scan("sensors")    // scan sensors table
  .window(Tumble over 1.hour on 'rowtime' as 'w') // define 1-hour window
  .groupBy('w, 'room)                          // group by window and room
  .select('room, 'w.end, 'temp.avg as 'avgTemp) // compute average temperature
```

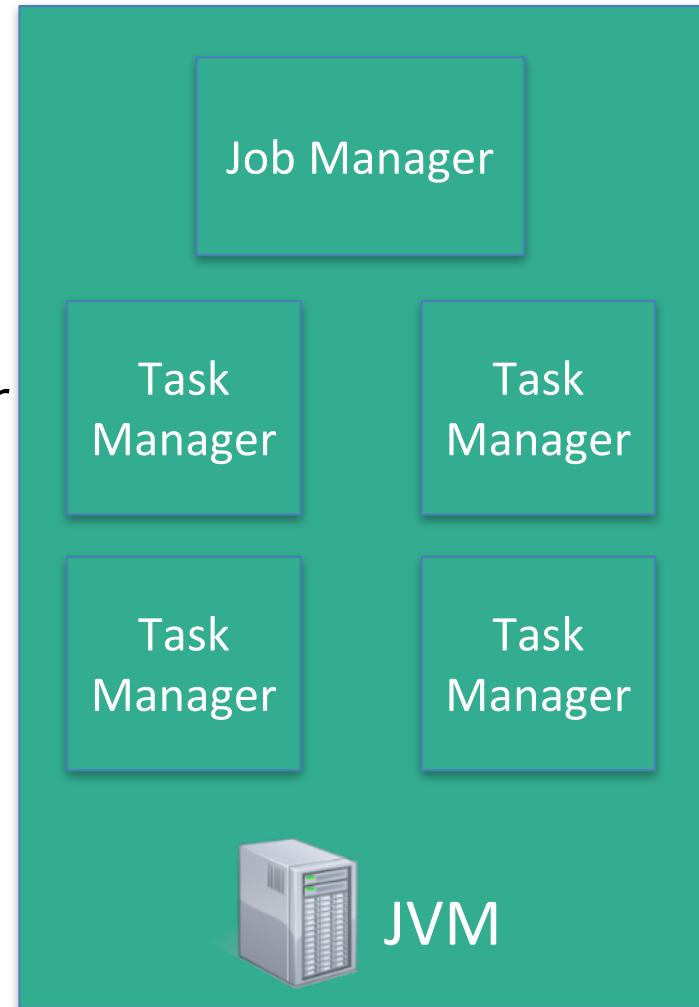
```
|SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp) AS avgTemp
|FROM sensors
|GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```

Deployment Options

Local Execution



- Starts local Flink cluster
- All processes run in the same JVM
- Behaves just like a regular Cluster
- Local cluster can be started in your IDE!
- Very useful for developing and debugging



Remote Execution



Submit
job

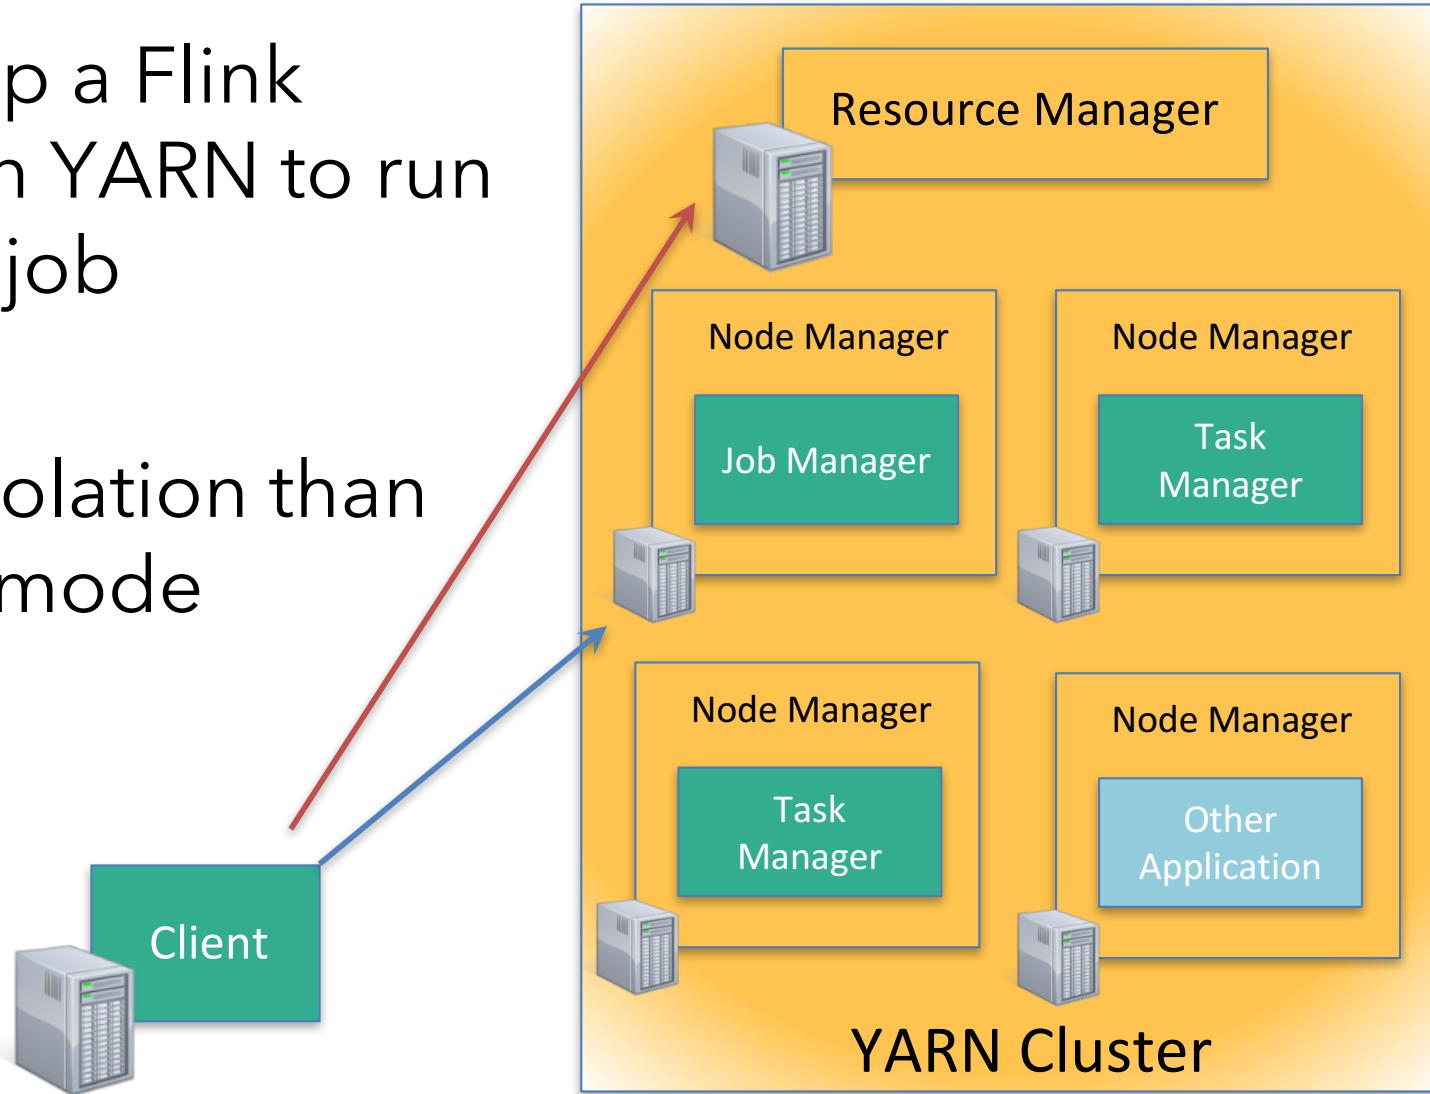


- **Submit** a Job to a remotely running cluster
- **Monitor** the status of a job

YARN Job Mode



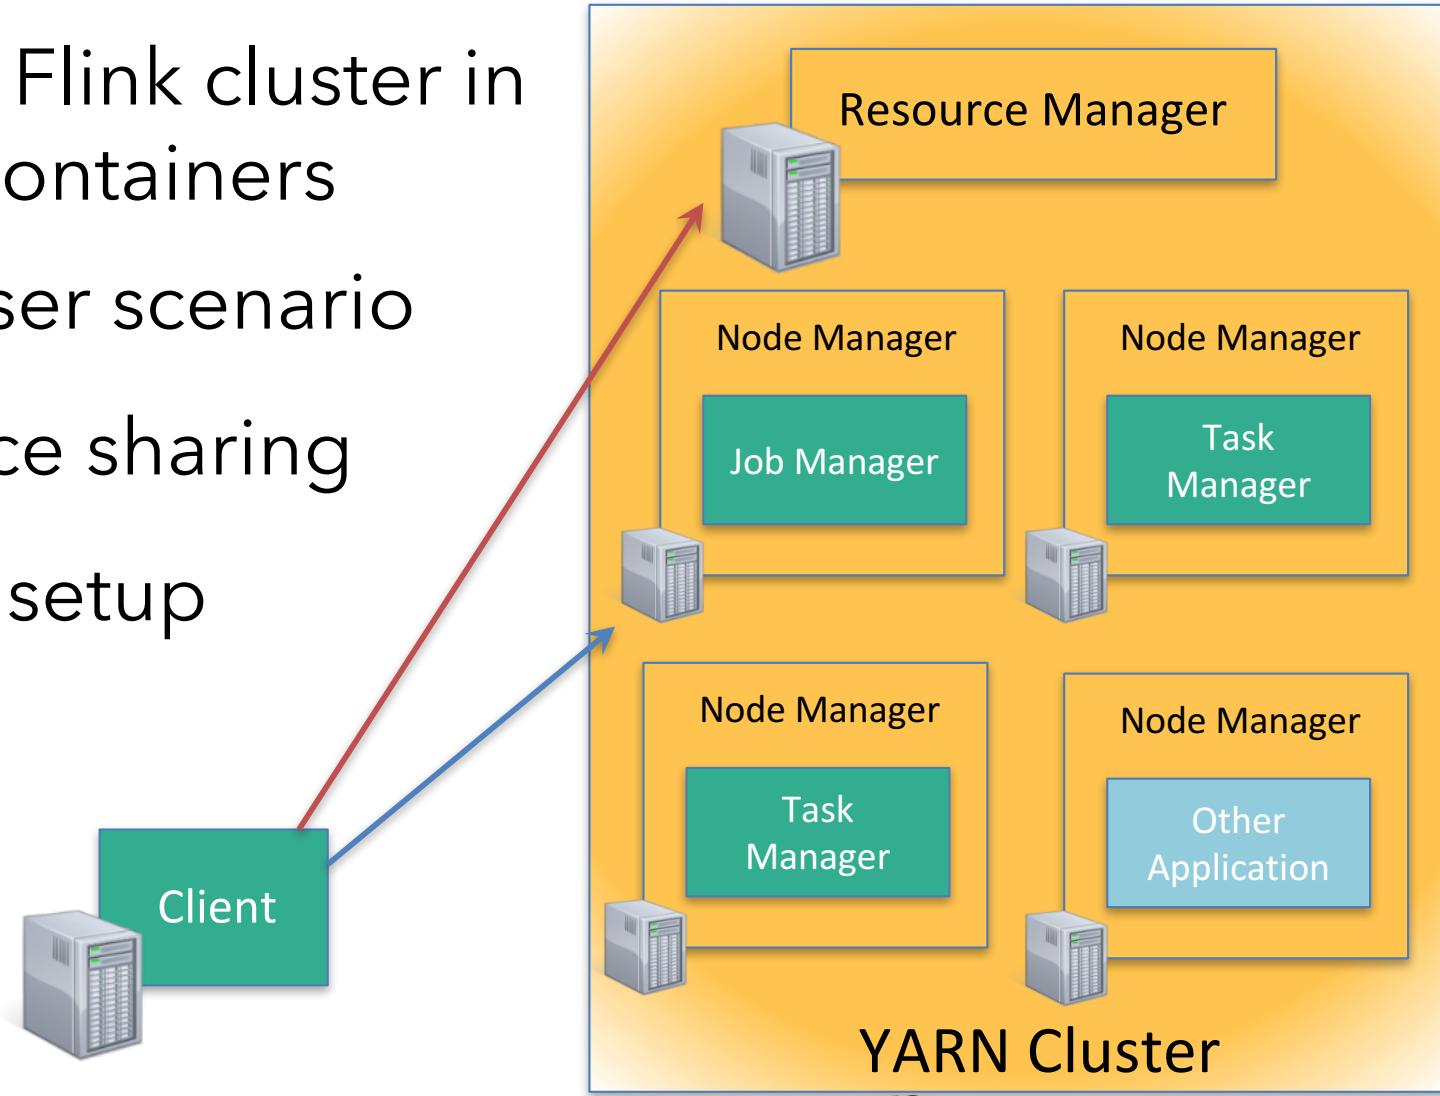
- Brings up a Flink cluster in YARN to run a single job
- Better isolation than session mode



YARN Session Mode



- Starts a Flink cluster in YARN containers
- Multi-user scenario
- Resource sharing
- Easy to setup





Other Deployment Options

- Apache Mesos
 - Either with or without DC/OS
- Amazon Elastic MapReduce
 - Available in EMR 5.1.0
- Google Compute Engine
 - Available via bdutil
- Docker / Kubernetes

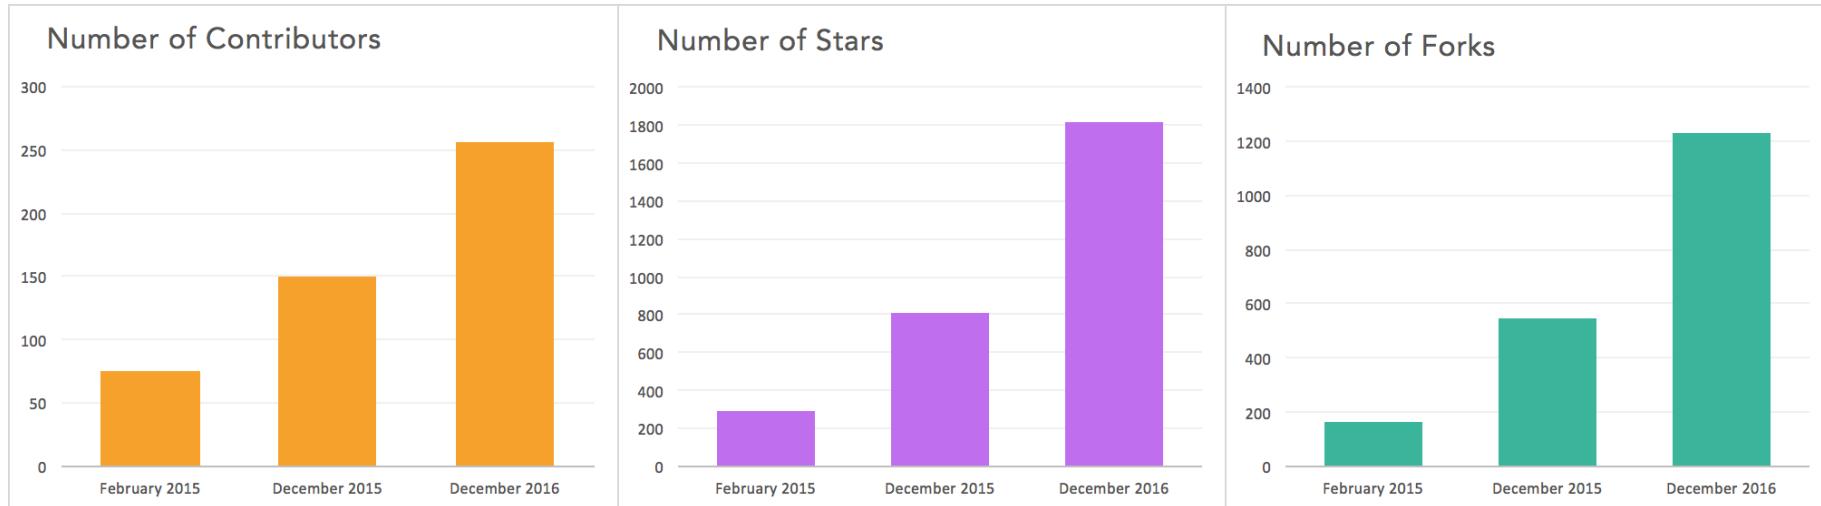


Flink in the real world

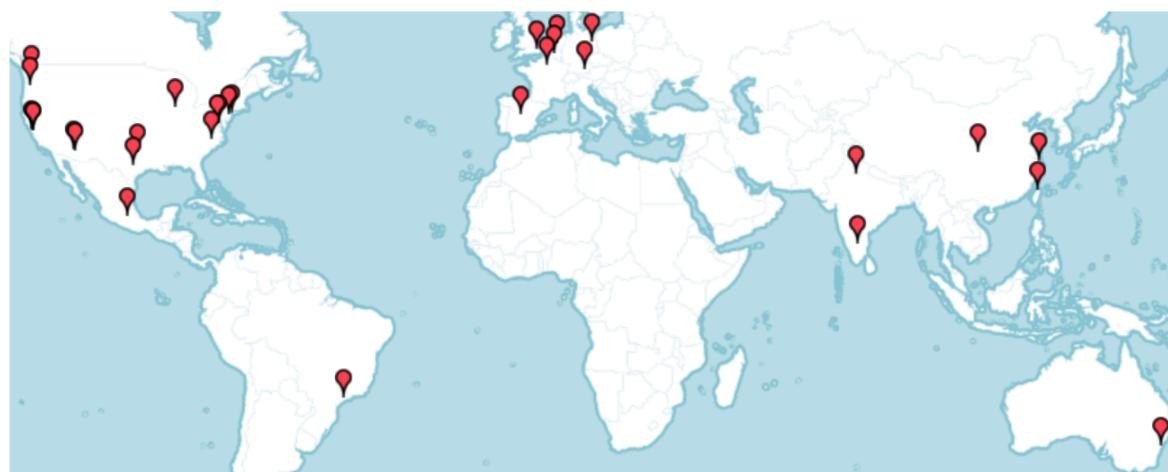


Flink community

Github



41 meetups
16,544 members



Powered by Flink



Zalando, one of the largest ecommerce companies in Europe, uses Flink for real-time business process monitoring.



King, the creators of Candy Crush Saga, uses Flink to provide data science teams with real-time analytics.



Alibaba, the world's largest retailer, built a Flink-based system (Blink) to optimize search rankings in real time.



Bouygues Telecom uses Flink for real-time event processing over billions of Kafka messages per day.

See more at flink.apache.org/poweredsby.html





Largest job has > 20 operators, runs on > 5000 vCores in 1000-node cluster, processes millions of events per second



Complex jobs of > 30 operators running 24/7, processing 30 billion events daily, maintaining state of 100s of GB with exactly-once guarantees



30 Flink applications in production for more than one year. 10 billion events (2TB) processed daily

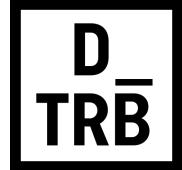
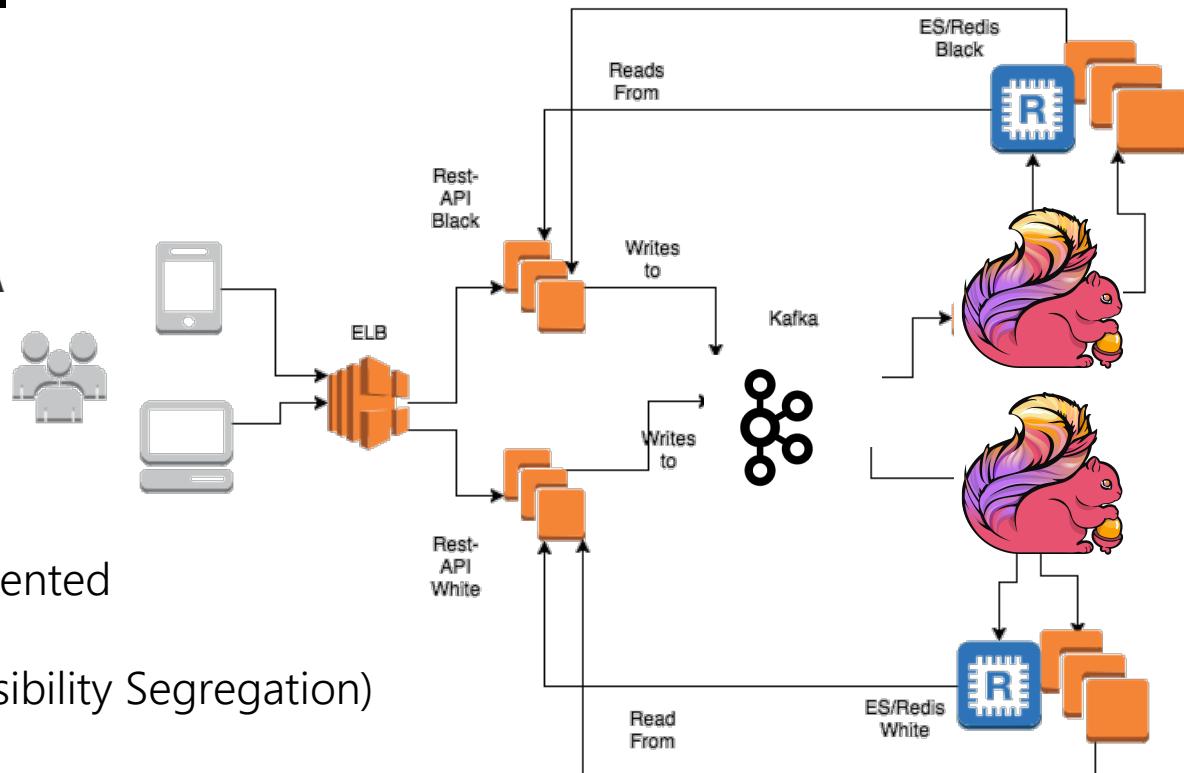


What is being built with Flink?

- First wave for streaming was **lambda architecture**
 - Aid batch systems to be more real-time
- Second wave was **analytics (real time and lag-time)**
 - Based on distributed collections, functions, and windows
- The next wave is **much broader:**
A new architecture for event-driven applications



@

THE SOCIAL NETWORK
FOR PETROLHEADS

Complete social network implemented
using event sourcing and
CQRS (Command Query Responsibility Segregation)



Flink Forward 2016

dataArtisans

EMC²

Google

King



otto group



redhat.



ResearchGate



cloudera

MAPR

TNG TECHNOLOGY
CONSULTING

SICS

MTA SZTAKI



mgm

people pattern

CRS4
IDEAS BECOME LIFE

ScaDS
DRESDEN LEIPZIG

Flink Forward 2017



San Francisco

- 10-11 April 2017
- The first Flink Forward event outside of Berlin
- Talks are online at sf.flink-forward.org/

Berlin

- 11-13 September 2017
- Over 350 attendees last year
- Registration opening soon!



<http://dataartisans.github.io/flink-training/>